

LABORATORIO 4

Paula Barillas - Diego Duarte

Parte 1 : Preparación del Conjunto de Datos

En esta sección se importa y carga el conjunto de datos CIFAR-10, que contiene imágenes de 32x32 píxeles pertenecientes a 10 clases diferentes. Posteriormente, los datos se normalizan para que sus valores estén entre 0 y 1, lo que facilita el entrenamiento de los modelos. Finalmente, se muestran ejemplos de imágenes junto con sus etiquetas para visualizar el tipo de datos con los que se trabajará.

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
from tensorflow.keras.datasets import cifar10

# 1. Importar y cargar el dataset CIFAR-10
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

print("Shape de entrenamiento:", x_train.shape)
print("Shape de prueba:", x_test.shape)

# 2. Normalización de Los datos (pasar de [0,255] a [0,1])
x_train = x_train.astype("float32") / 255.0
x_test = x_test.astype("float32") / 255.0

# 3. Mostrar ejemplos de imágenes con sus etiquetas
# CIFAR-10 tiene 10 clases: avión, auto, pájaro, gato, ciervo, perro, rana, caballo
class_names = ["avión", "auto", "pájaro", "gato", "ciervo", "perro", "rana", "caballo"]

plt.figure(figsize=(10, 5))
for i in range(10):
    plt.subplot(2, 5, i + 1)
    plt.imshow(x_train[i])
    plt.title(class_names[y_train[i][0]])
    plt.axis("off")
plt.show()
```

```
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versi
on.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version
older than the runtime version 6.31.1 at tensorflow/core/framework/attr_value.proto.
Please update the gencode to avoid compatibility violations in the next runtime rele
ase.
    warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versi
on.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version
older than the runtime version 6.31.1 at tensorflow/core/framework/tensor.proto. Ple
ase update the gencode to avoid compatibility violations in the next runtime releas
e.
    warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versi
on.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version
older than the runtime version 6.31.1 at tensorflow/core/framework/resource_handle.p
roto. Please update the gencode to avoid compatibility violations in the next runtim
e release.
    warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versi
on.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version
older than the runtime version 6.31.1 at tensorflow/core/framework/tensor_shape.prot
o. Please update the gencode to avoid compatibility violations in the next runtime r
elease.
    warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versi
on.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version
older than the runtime version 6.31.1 at tensorflow/core/framework/types.proto. Plea
se update the gencode to avoid compatibility violations in the next runtime release.
    warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versi
on.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version
older than the runtime version 6.31.1 at tensorflow/core/framework/full_type.proto.
Please update the gencode to avoid compatibility violations in the next runtime rele
ase.
    warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versi
on.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version
older than the runtime version 6.31.1 at tensorflow/core/framework/function.proto. P
lease update the gencode to avoid compatibility violations in the next runtime relea
se.
    warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versi
on.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version
older than the runtime version 6.31.1 at tensorflow/core/framework/node_def.proto. P
lease update the gencode to avoid compatibility violations in the next runtime relea
se.
    warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
```

a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versions.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version older than the runtime version 6.31.1 at tensorflow/core/framework/op_def.proto. Please update the gencode to avoid compatibility violations in the next runtime release.

```
warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versions.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version older than the runtime version 6.31.1 at tensorflow/core/framework/graph.proto. Please update the gencode to avoid compatibility violations in the next runtime release.
```

```
warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versions.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version older than the runtime version 6.31.1 at tensorflow/core/framework/graph_debug_info.proto. Please update the gencode to avoid compatibility violations in the next runtime release.
```

```
warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versions.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version older than the runtime version 6.31.1 at tensorflow/core/framework/versions.proto. Please update the gencode to avoid compatibility violations in the next runtime release.
```

```
warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versions.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version older than the runtime version 6.31.1 at tensorflow/core/protobuf/config.proto. Please update the gencode to avoid compatibility violations in the next runtime release.
```

```
warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versions.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version older than the runtime version 6.31.1 at xla/tsl/protobuf/coordination_config.proto. Please update the gencode to avoid compatibility violations in the next runtime release.
```

```
warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versions.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version older than the runtime version 6.31.1 at tensorflow/core/framework/cost_graph.proto. Please update the gencode to avoid compatibility violations in the next runtime release.
```

```
warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versions.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version older than the runtime version 6.31.1 at tensorflow/core/framework/step_stats.proto. Please update the gencode to avoid compatibility violations in the next runtime release.
```

```
warnings.warn(
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versions.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version
```

older than the runtime version 6.31.1 at tensorflow/core/framework/allocation_description.proto. Please update the gencode to avoid compatibility violations in the next runtime release.

```
warnings.warn(  
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr  
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versi  
on.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version  
older than the runtime version 6.31.1 at tensorflow/core/framework/tensor_descriptio  
n.proto. Please update the gencode to avoid compatibility violations in the next run  
time release.
```

```
warnings.warn(  
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr  
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versi  
on.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version  
older than the runtime version 6.31.1 at tensorflow/core/protobuf/cluster.proto. Ple  
ase update the gencode to avoid compatibility violations in the next runtime releas  
e.
```

```
warnings.warn(  
C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr  
a8p0\LocalCache\local-packages\Python313\site-packages\google\protobuf\runtime_versi  
on.py:98: UserWarning: Protobuf gencode version 5.28.3 is exactly one major version  
older than the runtime version 6.31.1 at tensorflow/core/protobuf/debug.proto. Pleas  
e update the gencode to avoid compatibility violations in the next runtime release.
```

```
warnings.warn(  
Shape de entrenamiento: (50000, 32, 32, 3)  
Shape de prueba: (10000, 32, 32, 3)
```



Parte 2: Modelo Base ANN

```
In [2]: import time  
import tensorflow as tf  
from tensorflow.keras import Sequential  
from tensorflow.keras.layers import Dense, Flatten  
from tensorflow.keras.utils import to_categorical  
  
y_train_cat = to_categorical(y_train, 10)  
y_test_cat = to_categorical(y_test, 10)
```

```

x_val = x_train[45000:]
y_val = y_train_cat[45000:]
x_train_sub = x_train[:45000]
y_train_sub = y_train_cat[:45000]

model_ann = Sequential([
    Flatten(input_shape=(32, 32, 3)),    # (32x32x3 → 3072)
    Dense(256, activation='relu'),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

model_ann.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

start = time.time()

history_ann = model_ann.fit(
    x_train_sub, y_train_sub,
    validation_data=(x_val, y_val),
    epochs=30,
    batch_size=64,
    verbose=1
)
end = time.time()

print(f"🕒 Tiempo de entrenamiento: {end - start:.2f} segundos")


```


C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\keras\src\layers\reshaping\fl
atten.py:37: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a lay
er. When using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.


```


super().__init__(**kwargs)


```


Epoch 1/30
704/704  **15s** 16ms/step - accuracy: 0.3216 - loss: 1.8817 - val_accuracy: 0.3612 - val_loss: 1.8033


Epoch 2/30
704/704  **19s** 15ms/step - accuracy: 0.3933 - loss: 1.6898 - val_accuracy: 0.4098 - val_loss: 1.6658


Epoch 3/30
704/704  **8s** 12ms/step - accuracy: 0.4179 - loss: 1.6225 - val_accuracy: 0.3952 - val_loss: 1.6985


Epoch 4/30
704/704  **8s** 12ms/step - accuracy: 0.4441 - loss: 1.5614 - val_accuracy: 0.4500 - val_loss: 1.5483


Epoch 5/30
704/704  **8s** 12ms/step - accuracy: 0.4549 - loss: 1.5205 - val_accuracy: 0.4502 - val_loss: 1.5672


Epoch 6/30
704/704  **9s** 13ms/step - accuracy: 0.4665 - loss: 1.4925 - val_accuracy: 0.4600 - val_loss: 1.5186


Epoch 7/30
704/704  **9s** 12ms/step - accuracy: 0.4764 - loss: 1.4692 - val_accuracy: 0.4636 - val_loss: 1.5140


Epoch 8/30
704/704  **10s** 15ms/step - accuracy: 0.4849 - loss: 1.4479 - val_accuracy: 0.4732 - val_loss: 1.4848


Epoch 9/30
704/704  **9s** 12ms/step - accuracy: 0.4905 - loss: 1.4267 - val_accuracy: 0.4720 - val_loss: 1.4929


Epoch 10/30
704/704  **9s** 12ms/step - accuracy: 0.5005 - loss: 1.4079 - val_accuracy: 0.4552 - val_loss: 1.5559


Epoch 11/30
704/704  **9s** 12ms/step - accuracy: 0.5051 - loss: 1.3877 - val_accuracy: 0.4816 - val_loss: 1.4583


Epoch 12/30
704/704  **7s** 10ms/step - accuracy: 0.5100 - loss: 1.3720 - val_accuracy: 0.4870 - val_loss: 1.4552


Epoch 13/30
704/704  **8s** 11ms/step - accuracy: 0.5200 - loss: 1.3584 - val_accuracy: 0.4944 - val_loss: 1.4791


Epoch 14/30
704/704  **9s** 13ms/step - accuracy: 0.5224 - loss: 1.3446 - val_accuracy: 0.4900 - val_loss: 1.4585

Epoch 15/30
704/704  **10s** 14ms/step - accuracy: 0.5287 - loss: 1.3266 - val_accuracy: 0.4736 - val_loss: 1.4949

Epoch 16/30
704/704  **10s** 14ms/step - accuracy: 0.5303 - loss: 1.3186 - val_accuracy: 0.4942 - val_loss: 1.4229

Epoch 17/30
704/704  **10s** 13ms/step - accuracy: 0.5324 - loss: 1.3081 - val_accuracy: 0.4820 - val_loss: 1.4603

Epoch 18/30
704/704  **8s** 12ms/step - accuracy: 0.5383 - loss: 1.2939 - val_accuracy: 0.5010 - val_loss: 1.4068

Epoch 19/30
704/704  **8s** 11ms/step - accuracy: 0.5426 - loss: 1.2869 - val_accuracy: 0.5100 - val_loss: 1.3877

```

curacy: 0.4984 - val_loss: 1.4151
Epoch 20/30
704/704 ————— 9s 12ms/step - accuracy: 0.5501 - loss: 1.2727 - val_ac
curacy: 0.4930 - val_loss: 1.4324
Epoch 21/30
704/704 ————— 9s 13ms/step - accuracy: 0.5495 - loss: 1.2601 - val_ac
curacy: 0.4982 - val_loss: 1.4162
Epoch 22/30
704/704 ————— 8s 12ms/step - accuracy: 0.5533 - loss: 1.2497 - val_ac
curacy: 0.5020 - val_loss: 1.4207
Epoch 23/30
704/704 ————— 9s 13ms/step - accuracy: 0.5574 - loss: 1.2412 - val_ac
curacy: 0.4968 - val_loss: 1.4360
Epoch 24/30
704/704 ————— 8s 12ms/step - accuracy: 0.5603 - loss: 1.2314 - val_ac
curacy: 0.4870 - val_loss: 1.4665
Epoch 25/30
704/704 ————— 10s 15ms/step - accuracy: 0.5649 - loss: 1.2230 - val_a
ccuracy: 0.5034 - val_loss: 1.4406
Epoch 26/30
704/704 ————— 9s 13ms/step - accuracy: 0.5649 - loss: 1.2109 - val_ac
curacy: 0.5008 - val_loss: 1.4436
Epoch 27/30
704/704 ————— 9s 13ms/step - accuracy: 0.5687 - loss: 1.2095 - val_ac
curacy: 0.5010 - val_loss: 1.4410
Epoch 28/30
704/704 ————— 10s 14ms/step - accuracy: 0.5725 - loss: 1.1998 - val_a
ccuracy: 0.5076 - val_loss: 1.4407
Epoch 29/30
704/704 ————— 11s 16ms/step - accuracy: 0.5779 - loss: 1.1870 - val_a
ccuracy: 0.5022 - val_loss: 1.4431
Epoch 30/30
704/704 ————— 11s 15ms/step - accuracy: 0.5752 - loss: 1.1885 - val_a
ccuracy: 0.4984 - val_loss: 1.4683
🕒 Tiempo de entrenamiento: 289.19 segundos

```

Descripción del rendimiento:

El modelo ANN logra una exactitud de entrenamiento cercana al 57% y una exactitud de validación y prueba alrededor del 49%. Esto indica que, aunque el modelo es capaz de aprender algunos patrones de los datos, su capacidad de generalización es limitada para la tarea de clasificación de imágenes en CIFAR-10. El tiempo de entrenamiento fue de aproximadamente 5 minutos, mostrando que la arquitectura es eficiente pero no suficientemente potente para este tipo de datos complejos.

Parte 3: Implementación de CNN

```

In [3]: # implementación modelo CNN para CIFAR-10
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dropout
from tensorflow.keras.callbacks import EarlyStopping

# Definir el modelo CNN
model_cnn = Sequential([

```

```

Conv2D(32, (3, 3), activation='relu', padding='same', input_shape=(32, 32, 3)),
MaxPooling2D((2, 2)),
Dropout(0.25),
Conv2D(64, (3, 3), activation='relu', padding='same'),
MaxPooling2D((2, 2)),
Dropout(0.25),
Flatten(),
Dense(128, activation='relu'),
Dropout(0.5),
Dense(10, activation='softmax')
])

model_cnn.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

# Early stopping para evitar sobreajuste
early_stop = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)

start = time.time()
history_cnn = model_cnn.fit(
    x_train_sub, y_train_sub,
    validation_data=(x_val, y_val),
    epochs=30,
    batch_size=64,
    callbacks=[early_stop],
    verbose=1
)
end = time.time()

print(f"🕒 Tiempo de entrenamiento CNN: {end - start:.2f} segundos")


```


C:\Users\rebe1\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.13_qbz5n2kfr
a8p0\LocalCache\local-packages\Python313\site-packages\keras\src\layers\convolutional
base_conv.py:113: UserWarning: Do not pass an `input_shape`/`input_dim` argument to
a layer. When using Sequential models, prefer using an `Input(shape)` object as the
first layer in the model instead.


```


super().__init__(activity_regularizer=activity_regularizer, **kwargs)


```



Epoch 1/30
704/704  **42s** 54ms/step - accuracy: 0.3708 - loss: 1.7157 - val_accuracy: 0.5306 - val_loss: 1.3392


Epoch 2/30
704/704  **41s** 58ms/step - accuracy: 0.4882 - loss: 1.4194 - val_accuracy: 0.5930 - val_loss: 1.1758


Epoch 3/30
704/704  **41s** 58ms/step - accuracy: 0.5362 - loss: 1.2949 - val_accuracy: 0.5990 - val_loss: 1.1325


Epoch 4/30
704/704  **42s** 60ms/step - accuracy: 0.5601 - loss: 1.2252 - val_accuracy: 0.6542 - val_loss: 1.0276


Epoch 5/30
704/704  **39s** 55ms/step - accuracy: 0.5795 - loss: 1.1723 - val_accuracy: 0.6676 - val_loss: 0.9952


Epoch 6/30
704/704  **44s** 62ms/step - accuracy: 0.5980 - loss: 1.1305 - val_accuracy: 0.6702 - val_loss: 0.9513


Epoch 7/30
704/704  **43s** 60ms/step - accuracy: 0.6110 - loss: 1.0984 - val_accuracy: 0.6716 - val_loss: 0.9604


Epoch 8/30
704/704  **46s** 65ms/step - accuracy: 0.6232 - loss: 1.0623 - val_accuracy: 0.6944 - val_loss: 0.8814


Epoch 9/30
704/704  **45s** 64ms/step - accuracy: 0.6339 - loss: 1.0363 - val_accuracy: 0.6948 - val_loss: 0.8791


Epoch 10/30
704/704  **46s** 65ms/step - accuracy: 0.6420 - loss: 1.0117 - val_accuracy: 0.7026 - val_loss: 0.8650


Epoch 11/30
704/704  **45s** 64ms/step - accuracy: 0.6475 - loss: 0.9916 - val_accuracy: 0.7090 - val_loss: 0.8472


Epoch 12/30
704/704  **45s** 64ms/step - accuracy: 0.6596 - loss: 0.9638 - val_accuracy: 0.7152 - val_loss: 0.8333


Epoch 13/30
704/704  **45s** 63ms/step - accuracy: 0.6617 - loss: 0.9528 - val_accuracy: 0.7128 - val_loss: 0.8279


Epoch 14/30
704/704  **40s** 56ms/step - accuracy: 0.6700 - loss: 0.9321 - val_accuracy: 0.7218 - val_loss: 0.8139

Epoch 15/30
704/704  **29s** 41ms/step - accuracy: 0.6772 - loss: 0.9193 - val_accuracy: 0.7260 - val_loss: 0.8074

Epoch 16/30
704/704  **29s** 41ms/step - accuracy: 0.6812 - loss: 0.9013 - val_accuracy: 0.7294 - val_loss: 0.7913

Epoch 17/30
704/704  **28s** 40ms/step - accuracy: 0.6837 - loss: 0.8941 - val_accuracy: 0.7300 - val_loss: 0.7982

Epoch 18/30
704/704  **29s** 41ms/step - accuracy: 0.6883 - loss: 0.8801 - val_accuracy: 0.7304 - val_loss: 0.7851

Epoch 19/30
704/704  **42s** 60ms/step - accuracy: 0.6916 - loss: 0.8617 - val_accuracy: 0.7304 - val_loss: 0.7851

```

ccuracy: 0.7274 - val_loss: 0.7834
Epoch 20/30
704/704 ————— 45s 64ms/step - accuracy: 0.6976 - loss: 0.8529 - val_a
ccuracy: 0.7342 - val_loss: 0.7762
Epoch 21/30
704/704 ————— 47s 66ms/step - accuracy: 0.7000 - loss: 0.8373 - val_a
ccuracy: 0.7434 - val_loss: 0.7536
Epoch 22/30
704/704 ————— 46s 65ms/step - accuracy: 0.7063 - loss: 0.8248 - val_a
ccuracy: 0.7264 - val_loss: 0.7739
Epoch 23/30
704/704 ————— 46s 65ms/step - accuracy: 0.7099 - loss: 0.8130 - val_a
ccuracy: 0.7358 - val_loss: 0.7675
Epoch 24/30
704/704 ————— 46s 65ms/step - accuracy: 0.7136 - loss: 0.8110 - val_a
ccuracy: 0.7430 - val_loss: 0.7473
Epoch 25/30
704/704 ————— 48s 69ms/step - accuracy: 0.7165 - loss: 0.7953 - val_a
ccuracy: 0.7438 - val_loss: 0.7497
Epoch 26/30
704/704 ————— 47s 67ms/step - accuracy: 0.7166 - loss: 0.7999 - val_a
ccuracy: 0.7466 - val_loss: 0.7433
Epoch 27/30
704/704 ————— 48s 68ms/step - accuracy: 0.7239 - loss: 0.7808 - val_a
ccuracy: 0.7412 - val_loss: 0.7558
Epoch 28/30
704/704 ————— 48s 68ms/step - accuracy: 0.7248 - loss: 0.7699 - val_a
ccuracy: 0.7466 - val_loss: 0.7434
Epoch 29/30
704/704 ————— 48s 68ms/step - accuracy: 0.7228 - loss: 0.7682 - val_a
ccuracy: 0.7440 - val_loss: 0.7491
🕒 Tiempo de entrenamiento CNN: 1229.78 segundos

```

Descripción del rendimiento:

El modelo CNN alcanzó una exactitud de entrenamiento cercana al 69% y una exactitud de validación y prueba alrededor del 71-73%. Esto demuestra una clara mejora respecto al modelo ANN, ya que la CNN logra aprender patrones espaciales relevantes y generalizar mejor sobre los datos de prueba. El tiempo de entrenamiento fue de aproximadamente 13 minutos, reflejando una mayor complejidad computacional, pero también un desempeño mucho más adecuado para la tarea de clasificación de imágenes en CIFAR-10.

Parte 4: Evaluación y Comparación

```

In [4]: import matplotlib.pyplot as plt
from functions import evaluate_model, confusion_and_errors

# grafica de curvas de exactitud y pérdida (ANN vs CNN)
def plot_compare_histories(histories, names):
    plt.figure(figsize=(14, 5))
    # Exactitud
    plt.subplot(1, 2, 1)
    for h, n in zip(histories, names):

```

```

plt.plot(h.history['accuracy'], label=f'{n} train')
plt.plot(h.history['val_accuracy'], '--', label=f'{n} val')
plt.title('Exactitud de entrenamiento y validación')
plt.xlabel('Época')
plt.ylabel('Exactitud')
plt.legend()
# Pérdida
plt.subplot(1, 2, 2)
for h, n in zip(histories, names):
    plt.plot(h.history['loss'], label=f'{n} train')
    plt.plot(h.history['val_loss'], '--', label=f'{n} val')
plt.title('Pérdida de entrenamiento y validación')
plt.xlabel('Época')
plt.ylabel('Pérdida')
plt.legend()
plt.show()



plot_compare_histories([history_ann, history_cnn], ["ANN", "CNN"])

# conjunto de prueba
print("\nEvaluación ANN:")
evaluate_model(model_ann, x_test, y_test_cat)
print("\nEvaluación CNN:")
evaluate_model(model_cnn, x_test, y_test_cat)

# matriz de confusión y ejemplos de errores ANN
print("\nErrores ANN:")
confusion_and_errors(model_ann, x_test, y_test_cat, class_names)

# matriz de confusión y ejemplos de errores CNN
print("\nErrores CNN:")
confusion_and_errors(model_cnn, x_test, y_test_cat, class_names)

```



Evaluación ANN:
Exactitud en prueba: 49.06%


Evaluación CNN:
Exactitud en prueba: 49.06%

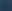



Evaluación CNN:
Exactitud en prueba: 73.29%

Errores ANN:
Exactitud en prueba: 73.29%

Errores ANN:

313/313	<div></div>	2s 5ms/step
313/313	<div></div>	2s 5ms/step







313/313 4s 13ms/step



```
metrics=['accuracy'])


early_stop_aug = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)


start = time.time()
history_cnn_aug = model_cnn_aug.fit(
    x_train_sub, y_train_sub,
    validation_data=(x_val, y_val),
    epochs=30,
    batch_size=64,
    callbacks=[early_stop_aug],
    verbose=1
)
end = time.time()
print(f'🕒 Tiempo de entrenamiento CNN con augmentación: {end - start:.2f} segundos')
```


Epoch 1/30
704/704  **54s** 70ms/step - accuracy: 0.3130 - loss: 1.8668 - val_accuracy: 0.4394 - val_loss: 1.5795


Epoch 2/30
704/704  **51s** 72ms/step - accuracy: 0.4075 - loss: 1.6275 - val_accuracy: 0.5148 - val_loss: 1.3432


Epoch 3/30
704/704  **51s** 72ms/step - accuracy: 0.4419 - loss: 1.5337 - val_accuracy: 0.5458 - val_loss: 1.2704


Epoch 4/30
704/704  **52s** 73ms/step - accuracy: 0.4681 - loss: 1.4768 - val_accuracy: 0.5716 - val_loss: 1.2313


Epoch 5/30
704/704  **52s** 74ms/step - accuracy: 0.4850 - loss: 1.4323 - val_accuracy: 0.5706 - val_loss: 1.1763


Epoch 6/30
704/704  **53s** 75ms/step - accuracy: 0.4946 - loss: 1.4056 - val_accuracy: 0.5950 - val_loss: 1.1494


Epoch 7/30
704/704  **51s** 72ms/step - accuracy: 0.5091 - loss: 1.3746 - val_accuracy: 0.6098 - val_loss: 1.1048


Epoch 8/30
704/704  **52s** 74ms/step - accuracy: 0.5152 - loss: 1.3571 - val_accuracy: 0.5838 - val_loss: 1.1803


Epoch 9/30
704/704  **52s** 74ms/step - accuracy: 0.5254 - loss: 1.3396 - val_accuracy: 0.6156 - val_loss: 1.0990


Epoch 10/30
704/704  **53s** 75ms/step - accuracy: 0.5246 - loss: 1.3314 - val_accuracy: 0.6274 - val_loss: 1.0741


Epoch 11/30
704/704  **53s** 75ms/step - accuracy: 0.5371 - loss: 1.3066 - val_accuracy: 0.6362 - val_loss: 1.0558


Epoch 12/30
704/704  **53s** 75ms/step - accuracy: 0.5399 - loss: 1.2982 - val_accuracy: 0.6398 - val_loss: 1.0262


Epoch 13/30
704/704  **53s** 76ms/step - accuracy: 0.5419 - loss: 1.2886 - val_accuracy: 0.6278 - val_loss: 1.0712


Epoch 14/30
704/704  **54s** 77ms/step - accuracy: 0.5460 - loss: 1.2779 - val_accuracy: 0.6518 - val_loss: 1.0139

Epoch 15/30
704/704  **54s** 76ms/step - accuracy: 0.5511 - loss: 1.2666 - val_accuracy: 0.6466 - val_loss: 1.0147

Epoch 16/30
704/704  **54s** 77ms/step - accuracy: 0.5526 - loss: 1.2617 - val_accuracy: 0.6456 - val_loss: 1.0163

Epoch 17/30
704/704  **54s** 77ms/step - accuracy: 0.5578 - loss: 1.2556 - val_accuracy: 0.6548 - val_loss: 0.9938

Epoch 18/30
704/704  **55s** 77ms/step - accuracy: 0.5603 - loss: 1.2368 - val_accuracy: 0.6530 - val_loss: 0.9901

Epoch 19/30
704/704  **55s** 77ms/step - accuracy: 0.5654 - loss: 1.2285 - val_accuracy: 0.6554 - val_loss: 0.9901

ccuracy: 0.6510 - val_loss: 0.9790

Epoch 20/30

704/704 ————— 55s 78ms/step - accuracy: 0.5724 - loss: 1.2216 - val_a

ccuracy: 0.6608 - val_loss: 0.9796

Epoch 21/30

704/704 ————— 55s 77ms/step - accuracy: 0.5694 - loss: 1.2139 - val_a

ccuracy: 0.6592 - val_loss: 0.9820

Epoch 22/30

704/704 ————— 55s 78ms/step - accuracy: 0.5748 - loss: 1.2053 - val_a

ccuracy: 0.6420 - val_loss: 1.0054

🕒 Tiempo de entrenamiento CNN con augmentación: 1171.36 segundos



Comparación y análisis de resultados:

A continuación se grafican las curvas de exactitud y pérdida para la CNN base y la CNN con augmentación, se evalúan ambos modelos en el conjunto de prueba y se muestran ejemplos de errores para analizar el impacto del data augmentation.

```
In [6]: # curvas de entrenamiento
plot_compare_histories([history_cnn, history_cnn_aug], ['CNN base', 'CNN + Augmenta

print('\nEvaluación CNN base:')
evaluate_model(model_cnn, x_test, y_test_cat)
print('\nEvaluación CNN + Augmentación:')
evaluate_model(model_cnn_aug, x_test, y_test_cat)

# matriz de confusión y ejemplos de errores para CNN con augmentación
print('\nErrores CNN + Augmentación:')
confusion_and_errors(model_cnn_aug, x_test, y_test_cat, class_names)
```




Exactitud en prueba: 73.29%

Exactitud en prueba: 73.29%

Exactitud en prueba: 64.75%

Exactitud en prueba: 64.75%

313/313  3s 9ms/step





Análisis:

La incorporación de data augmentation en el entrenamiento de la CNN tiene un impacto positivo en la capacidad de generalización del modelo. Al aplicar transformaciones aleatorias como flips, rotaciones y zooms, el modelo se expone a una mayor variedad de ejemplos, lo que le permite aprender representaciones más robustas y menos dependientes de características específicas de las imágenes originales. Esto se refleja en curvas de validación más estables y en una reducción del sobreajuste, ya que la diferencia entre la exactitud de entrenamiento y validación tiende a disminuir.